

Optimizing Job Shop Scheduling in a Manufacturing Plant using Genetic Algorithm

Fawaz Khmous (1241518)

Amir Dar-Almasri (1240890)

Instructor: Dr. Yazan Abu-Farha

Department of Electrical and Computer Engineering
Birzeit University

Artificial Intelligence (ENCS3340) — Summer Course, 2026

1. Problem Formulation

1.1 Problem Description

The **Job Shop Scheduling Problem (JSSP)** is a prominent benchmark in combinatorial optimization and operations research. In modern manufacturing plants, discrete physical stations (such as cutting benches, precision drilling centers, and assembly bays) operate concurrently to process diverse goods.

Formally, a set of n independent jobs $\{J_1, J_2, \dots, J_n\}$ must be processed on a set of m dedicated machines $\{M_1, M_2, \dots, M_m\}$. Each job J_i follows a predetermined, technological sequence of operations:

$$J_i : O_{i,1} \rightarrow O_{i,2} \rightarrow \dots \rightarrow O_{i,k_i}$$

where operation $O_{i,k}$ requires exclusive processing on machine $M(O_{i,k})$ for an uninterrupted duration $d_{i,k} \geq 0$. The ordering of visited machines varies per job, inducing dense combinatorial dependencies.

1.2 Decision Variables

For each operation $O_{i,k}$, the scheduler determines:

- $\text{start}(O_{i,k}) \geq 0$: the time execution begins on the designated machine.
- $\text{end}(O_{i,k}) = \text{start}(O_{i,k}) + d_{i,k}$: completion time.

The output is an integrated plant-wide timeline assigning start and finish times to every operation without overlaps.

1.3 Hard Physical Constraints

1. Machine Non-Preemption (Capacity Constraint). Each machine processes at most one operation at any moment. Operations cannot be preempted once initiated. For any operations O_a, O_b with $M(O_a) = M(O_b)$:

$$\text{end}(O_a) \leq \text{start}(O_b) \vee \text{end}(O_b) \leq \text{start}(O_a)$$

2. Job Operation Precedence (Routing Constraint). Within any job J_i , operation $O_{i,k+1}$ cannot commence until $O_{i,k}$ has completely finished:

$$\text{start}(O_{i,k+1}) \geq \text{end}(O_{i,k}), \quad \forall i, k$$

1.4 Objective Function

The objective is to minimize the **Makespan** ($C_{\max}$), representing total factory completion time:

$$\min C_{\max} = \min \left(\max_{i,k} \text{end}(O_{i,k}) \right)$$

Minimizing makespan directly maximizes production throughput, lowers work-in-progress inventory, and reduces operating costs.

1.5 System Input & Output Format

Input Format. The system accepts job configuration files following the course specification:

```
Job_1:  M1[10] -> M2[5] -> M4[12]
Job_2:  M2[7] -> M3[15] -> M1[8]
```

The parser automatically extracts the machine routing, duration parameters, and job counts, adapting to arbitrary plant configurations.

Output Deliverables. The solver outputs:

- Detailed scheduling reports with machine utilization metrics.
- Chronological start/end tables for all operations.
- High-resolution color-coded Gantt charts.
- Evolutionary fitness convergence plots over generations.

2. Genetic Algorithm Design

2.1 Chromosome Representation

We utilize the academic **operation-based permutation encoding**. For a factory instance with n jobs where job j comprises k_j operations, the chromosome is represented as an integer vector of length $L = \sum_{j=1}^n k_j$. The identifier for job j appears exactly k_j times in the gene sequence.

The r -th occurrence of job ID j designates the r -th operation $O_{j,r}$ of job j .

Example. For two jobs with three operations each:

Genes: [1, 2, 1, 2, 1, 2]

The first 1 corresponds to $O_{1,1}$, the second 1 to $O_{1,2}$, and the third 1 to $O_{1,3}$.

Key Invariant: Any arbitrary permutation of this multiset is **intrinsically feasible** with respect to job precedence constraints, eliminating the need for complex repair heuristics.

2.2 Semi-Active Schedule Decoder

The chromosome genotype is decoded into an active physical schedule using an $\mathcal{O}(L)$ linear-time deterministic decoder. Two arrays monitor factory state:

- **jobEnd[j]**: completion time of the latest operation of job j .
- **machEnd[m]**: availability timestamp of machine m .

Iterating through the chromosome, each gene's corresponding operation is scheduled at the earliest valid moment:

$$\begin{aligned} \text{start} &= \max(\text{jobEnd}[j], \text{machEnd}[m]) \\ \text{end} &= \text{start} + d \\ C_{\max} &= \max(C_{\max}, \text{end}) \end{aligned}$$

Both tracking arrays are updated in place. This guarantees zero machine collisions and zero precedence violations by construction.

2.3 Crossover Operator: POX

Standard crossover operators create invalid operation frequencies. We implement **Precedence Operation Crossover (POX)**:

1. A non-empty subset of job IDs $J_{\text{keep}} \subset \{1, \dots, n\}$ is randomly selected.
2. In Child 1, operations belonging to J_{keep} retain their exact positional indices from Parent 1.
3. Remaining vacant positions are filled with operations not in J_{keep} in the order they appear in Parent 2.
4. Child 2 is generated symmetrically by inverting parent roles.

POX guarantees offspring validity, preserves desirable sub-schedules, and maintains population feasibility across generations.

2.4 Mutation Operator: Swap Mutation

To prevent premature stagnation, **swap mutation** picks two gene positions holding different job IDs and swaps them. This perturbs machine priority orderings while strictly conserving the multiset count. Parents are copied before mutation to preserve elite genomes.

2.5 Parent Selection & Linear Elitism

Tournament Selection ($k = 3$) randomly samples three individuals and selects the one with the lowest makespan, balancing selective pressure with diversity.

Additionally, **Elitism** ($E = 2$) carries the top two fittest chromosomes directly into the next generation using `heapq.nsmallest`, guaranteeing that best-so-far solutions are never degraded.

3. Experimental Setup & Benchmarks

3.1 Layered System Architecture

The system is structured in a strict Directed Acyclic Graph (DAG) with zero circular dependencies across three core modules:

- **Domain Layer** (`core/models.py`): Memory-efficient dataclasses (`slots=True`) for factory entities.
- **Genotype Layer** (`core/chromosome.py`): Encapsulated DNA, evaluation engine, and timeline builder.
- **Algorithm Layer** (`core/ga.py`): Genetic algorithm solver managing populations, operators, and elitism.
- **I/O & Visualization** (`io_handler/`, `visualization/`): Input parser, tabular reporter, and charting suites.

3.2 Hyperparameter Configuration

Table 1 presents the default parameters established through empirical experimentation.

Table 1: Default GA hyperparameter configuration.

Parameter	Function	Value
Population Size (P)	Solution candidate pool	100
Generations (G)	Iterative evolution budget	200
Crossover Rate (p_c)	POX recombination probability	0.85
Mutation Rate (p_m)	Swap operator probability	0.15
Tournament Size (k)	Selection pressure magnitude	3
Elitism Count (E)	Best individuals preserved	2

3.3 Empirical Test Case Evaluations

Test Case 1: Specification Example (2 Jobs $\times$ 4 Machines). The sample instance from the course specification:

```
Job_1:  M1[10] -> M2[5] -> M4[12]
Job_2:  M2[7] -> M3[15] -> M1[8]
```

Result: The solver achieved the **exact global optimum** makespan of **30** time units in 0.011 seconds.

Test Case 2: Factory Benchmark (10 Jobs $\times$ 10 Machines). A standard 100-operation industrial instance was parsed from `data/input.txt`. The solver reduced makespan by **24.51%** from 1,163 down to **878** time units (Table 2).

Table 2: Performance on the 10 $\times$ 10 manufacturing benchmark.

Metric	Observed Value
Total Factory Operations	100 operations
Initial Population Best Makespan	1,163 time units
Final Optimized Makespan ($C_{\max}$)	878 time units
Makespan Reduction	−285 units (24.51%)
Mean Machine Utilization	57.89%
Constraint Violations	0 (100% Feasible)

Test Case 3: Industrial Stress Test (30 Jobs $\times$ 20 Machines). An extreme instance with 600 operations (search space $\approx 10^{970}$) was tested with population 250 across 300 generations (75,000 schedule evaluations, 45 million gene operations).

Table 3: Industrial stress test benchmark results.

Metric	Observed Value
Problem Scale	30 Jobs $\times$ 20 Machines (600 ops)
Total Schedules Decoded	75,000 evaluations
Wall-Clock Execution Time	1.45 seconds
Evaluation Throughput	3,518,472 gene-ops/s
Initial $\rightarrow$ Final Makespan	3,824 $\rightarrow$ 2,549 (−33.34%)
Physical Violations	0

4. Visualizations & Comprehensive Analysis

4.1 Gantt Chart Timeline Schedule

Figure 1 illustrates the final optimized schedule for the 10-job $\times$ 10-machine plant ($C_{\max} = 878$). Machine rows (M_1 – M_{10}) depict ordered operation intervals with color-coded job identifiers. Hatched red blocks indicate unavoidable machine idle gaps caused by upstream job precedence constraints.




Figure 1: Optimized Gantt chart for the 10 $\times$ 10 plant ($C_{\max} = 878$).

4.2 Evolutionary Convergence Trajectory

Figure 2 shows fitness progression over 200 generations. The solid curve tracks the best makespan, exhibiting strict monotonic decrease via elitism. The dashed curve shows the population average makespan, confirming steady genetic convergence without diversity loss.




Figure 2: GA convergence trajectory: Best vs. average makespan.

4.3 Machine Utilization Breakdown

Table 4 details active machine utilization for the 10 $\times$ 10 schedule. Utilization is computed as the ratio of cumulative active processing time to total makespan:

$$U_m = \frac{\sum_{O \in \text{Mach}(m)} \text{duration}(O)}{C_{\max}} \times 100\%$$

Machine M_3 represents the plant bottleneck (68.68% active utilization), whereas M_8 experiences the greatest idle time (48.18%) due to precedence wait times.

Table 4: Machine utilization breakdown for the 10 $\times$ 10 plant.

Machine	Active Util. (%)	Machine	Active Util. (%)
M_1	55.47	M_6	67.88
M_2	65.38	M_7	60.14
M_3	68.68 (Bottleneck)	M_8	48.18
M_4	48.41	M_9	49.43
M_5	49.20	M_{10}	66.17
Factory Average		57.89%	

4.4 Automated Validation Suite

An automated test suite containing **72 exhaustive unit and integration tests** across 12 categories verified all system requirements with a **100% pass rate**:

- **Collision Checking:** Zero temporal overlaps on any machine across thousands of generated schedules.
- **Precedence Integrity:** Verified $\text{start}(O_{i,k+1}) \geq \text{end}(O_{i,k})$ across all job sequences.
- **Boundary Cases:** Successfully validated single-job, single-machine, and zero-duration scenarios.

4.5 Sensitivity Analysis & Demo Readiness

Parameter sweep experiments revealed that:

- Crossover rates in $[0.80, 0.90]$ achieve optimal schema exchange.
- Mutation rates below 0.05 cause premature stagnation, whereas rates above 0.25 disrupt convergence.
- Tournament size $k = 3$ reliably balances selective pressure with diversity.

4.6 Conclusions

The implemented Genetic Algorithm successfully solves the Job Shop Scheduling Problem by combining operation-based permutation encoding with a linear-time semi-active decoder. The system guarantees feasibility by construction, achieves significant makespan reductions across problem scales, and provides rich interactive visualizations.